

Policy Gradient Methods for Reinforcement Learning with Function Approximation

Reminders

- REINFORCE algorithm performs Monte Carlo sampling of the stochastic policy, and updates the policy using policy gradient calculated from the partial returns (G_t) of episodes
 - Unbiased estimate of the gradient, but slow convergence since it uses rollouts instead of function approximation
 - Improves greatly aided with value function approximation

Overview

- More focused on discrete actions (examples with softmax)
- Explicitly represent the policy with function approximator (parametrised)
- Update the parameters according to the gradient of expected reward
- Previous work:
 - Approximation of value-function + greedy *deterministic* policy
 - Optimal policy often stochastic and policy highly noise (from v-f) dependent
- Function approximators: e.g. NN with state as input and action probabilities as output
 - Update parameters using the policy gradient: $\Delta\theta = \alpha \frac{\partial J}{\partial \theta}$
- Here, small changes in θ cause only small changes in (stochastic) policy and thus state-visitation distribution, while small changes in v-f with deterministic policy have a larger influence
- Then, two different objectives used:
 - Global average reward:

$$J(\pi) = \lim_{n \rightarrow \infty} E_{\pi} \left[\sum_{i=1}^n r_i \right] = \sum_s d^{\pi}(s) \sum_a \pi(s, a | \theta) E[r | a, s]$$

- Long-term reward from a start state (**prevalent**):

$$J(\pi) = E_{\pi} \left[\sum_{t=1}^{\infty} \gamma^{t-1} r_t \mid s_0 \right]$$

Key ingredients

- Convergence of the policy iteration with arbitrary differentiable function approximation (local optimum)
- Policy gradient theorem:

$$\frac{\partial J}{\partial \theta} = \sum_s d^{\pi}(s) \sum_a \frac{\partial \pi(s, a | \theta)}{\partial \theta} Q^{\pi}(s, a)$$

- The effect of policy changes on the state distribution $\frac{\partial d^{\pi}(s)}{\partial \theta}$ does not appear
 - Convenient for sampling: simply draw samples following π to obtain unbiased estimate $\frac{\partial \log \pi(s, a | \theta)}{\partial \theta} Q^{\pi}(s, a)$ of the gradient
- Good, but still need to approximate Q!
 - Using actual returns leads to REINFORCE
 - Function approximation for Q^{π} speeds up learning and gives better performance
 - “Convenient” function approximation $Q^{\omega} = \frac{\partial \log \pi(s, a | \theta)}{\partial \theta}^T \omega$ (linear in ω) which we estimate from an unbiased Q^{π} estimate, e.g. the returns

Comments

- Well written and easy to follow

- Gives a lot of insights and useful comments on the matter
- Especially found useful comparisons with REINFORCE algorithm and insights about sampling
- Can see strong collaboration influences
- Sutton's book acts as a great supplement for understanding the paper